

Root Management	Delegated Votes	“Nuclear” Option	
	<p><b>Root Leader</b></p> <ul style="list-style-type: none"><li>• Broadcast command to all replicas.</li><li>• Resolve conflicting delegations by comparing subquorum terms.</li><li>• If current vote count is a majority, begin epoch transition.</li></ul> <p><b>Root Delegates</b></p> <ul style="list-style-type: none"><li>• if epoch &lt; current epoch: send no votes</li><li>• if vote undelegated: send self vote</li><li>• if candidate: send self vote</li><li>• if delegate: send all votes</li></ul> <p>Vote: (epoch <b>e</b>, quorum <b>q</b>, term <b>t</b>, votes <b>v</b>)</p>	<p>Delegations are only valid for the next epoch change. If enough delegates have failed that the epoch change cannot be made, a “nuclear” option resets delegates.</p> <p>Triggered by a nuclear timeout ≥ root election timeout to ensure root leader is dead and delegates can’t establish leader.</p> <ul style="list-style-type: none"><li>• Increment epoch beyond vote delegation limit, resetting all delegations.</li><li>• Conduct new root election/epoch change with all available replicas.</li><li>• Update health of all failed nodes and reconfigure epoch.</li></ul>	
Epoch Changes		Fuzzy Transitions	Epoch Decisions
<p>Initiated by request, reconfiguration, localization, quiescence procedures.</p> <p><b>Root Leader</b></p> <ul style="list-style-type: none"><li>• Monotonically increase epoch number, Define subquorums, initial leaders.</li><li>• Initiate delegated vote on epoch-change.</li><li>• On commit, begin fuzzy transition.</li></ul> <p><b>Subquorum Replicas</b></p> <ul style="list-style-type: none"><li>• Write tombstone into current log.</li><li>• Finalize commit for accesses prior to the tombstone record, forward new requests.</li><li>• On tombstone commit: truncate and archive log, join new subquorum configuration.</li></ul>		<p><b>Initiating:</b> leader of subquorum in e-1</p> <p><b>Remote:</b> leader of subquorum in e</p> <ul style="list-style-type: none"><li>• Initiating sends last committed command for every object required by remote, Null for objects without accesses, and number of outstanding entries.</li><li>• Remote appends last entries and performs batch consensus to bring subquorum to the Same state.</li><li>• On remote commit, reports to root leader and begins accepting new accesses.</li></ul> <p><b>Note:</b> background anti-entropy optimizes handoff process by reducing data volume.</p>	
Operations	Consensus and Accesses	Remote Accesses	
	<p>Clients are forwarded to the subquorum leader with responsibility for requested object(s).</p> <ul style="list-style-type: none"><li>• <b>Read(o):</b> Leader responds with last committed entry; marks response if uncommitted entry for object exists. Adds read access to log but does not begin consensus (aggregates reads with writes).</li><li>• <b>Write(o):</b> Leader increments objects version number and creates a corresponding log entry. Sends consensus request and responds to client when the entry is committed.</li></ul>	<p>In a multi-object transaction, remote accesses serialize inter-quorum access.</p> <p><b>Initiating:</b> append entries in log and send remote access request to remote leader.</p> <p><b>Remote:</b> create sub-epoch to demarcate remote access, add entry and respond to initiating replica when committed.</p> <p><b>Initiating:</b> on remote commit, create local sub-epoch, and commit entries appended to logs.</p>	